



RESEARCH EDITION • 2026

The Prompting Menu

Essential theories, techniques, examples
and safety checks for better AI prompts.

BEGINNER FRIENDLY

PRO TIPS INCLUDED



Build clearer prompts, layer by layer.

A practical reference for text, images, code and
real work.

Welcome to the menu

Prompting is how you give an AI a job, the information it needs and a clear shape for the answer.

This guide collects the **essential, evidence-backed prompting theories and techniques** most useful to everyday people and working professionals. It is not a list of magic words.

HOW TO READ EACH MENU ITEM

Theory explains it. **Use cases** show when it fits. **Example** demonstrates it. **Counter Tip** gives one action. **Pro Tip** adds depth. **Watch out** names the limit.

- 01 Start with a clear job.
- 02 Add only useful context.
- 03 Make success checkable.
- 04 Test the real result.
- 05 Review the risks.

Included

The seven-layer method, 24 practical techniques, text/image/code prompting, evaluation, BITE, risks, governance and professional workflow topics.

Not claimed

An exhaustive canon, universal model behaviour, legal advice, provider-specific guarantees or proof that an output is safe.

Research surveys describe dozens of text and multimodal techniques. Terminology and model behaviour continue to change. Examples here are illustrative and must be tested. [\[17\]](#)

Menu contents

Build your burger

- 04** Choose your route
- 05** Seven-layer burger
- 09** Worked example
- 10** Technique selector

Prompting Specials

- 11** Direct prompting
- 12** Role and audience
- 13** One-shot and few-shot
- 14** Structure and delimiters
- 15** Constraints
- 16** Decomposition and chaining
- 17** Worked reasoning
- 18** Self-consistency

Kitchen Tools + Special Orders

- 19** Retrieval and grounding
- 20** Tool use
- 21** Structured outputs
- 22** Multimodal prompting
- 23** Image prompting
- 24** Code prompting
- 25** Iterative prompting
- 26** Prompt evaluation

Professional Counter

- 27** Instruction hierarchy
- 28** Long context
- 29** Templates and versioning
- 30** Alternative paths

Professional Counter continued

- 31** Provider settings
- 32** Multilingual prompting
- 33** Cost and latency
- 34** Rights and provenance

Review and safety

- 35** BITE
- 36** Responsible-AI review
- 37** Hallucination
- 38** Prompt injection
- 39** Bias, privacy and judgement

Take away

- 40** Use-case menu
- 41** Complete workflow
- 42** Glossary
- 43** Sources
- 45** Final checklist

Choose your route

You do not need every technique. Choose the smallest route that solves the real problem.

WRITE OR SUMMARISE

Build the seven layers

Start with Goal and Task. Add source material, requirements and the output shape.

Then try: Direct prompting → examples → iteration

REPEAT A TASK

Engineer and evaluate

Define success, assemble representative test cases and compare prompt versions.

Then try: Few-shot → structured output → evaluation

MAKE AN IMAGE

Brief the visual

Name the use, subject, composition, look and what must not change.

Then try: Generate → inspect → edit one issue

BUILD OR FIX CODE

Brief the environment

Name versions, files, expected behaviour, boundaries and verification.

Then try: Plan → implement → test → review

BEFORE YOU PASTE ANYTHING

Remove secrets and unnecessary personal data. Check your chosen tool's current retention, training and access settings.

The seven-layer burger



Seven layers: Goal → Task → Context and Input → Requirements and Details → Style and Quality → Output Format → Rules and Boundaries.

Each ingredient has one job.
Stack only what the task needs.

1. **Goal** — why you need this.
2. **Task** — the exact job.
3. **Context and Input** — useful background and material.
4. **Requirements and Details** — what the result must include.
5. **Style and Quality** — how it should sound or look.
6. **Output Format** — the shape of the answer.
7. **Rules and Boundaries** — what must not happen.

THE WRAPPER IS NOT A SAFETY SWITCH

Rules and Boundaries state limits in the prompt. The responsible-AI review checks what could still go wrong.

Goal + Task

TOP BUN · REQUIRED

1. Goal

Theory: Why you need this, and what the result should help you achieve.

Ask: What should become possible after this is done?

Help new staff understand our returns policy.

TIP

Name the decision, action or outcome—not only the topic.

PATTY · REQUIRED

2. Task

Theory: The exact job you want the AI to do.

Ask: What action should the AI perform?

Summarise the policy and extract the five steps.

TIP

Lead with a strong verb: compare, extract, draft, classify or debug.

PRO TIP · GOAL IS NOT TASK

“Improve onboarding” is a Goal. “Write a one-page guide” is a Task. A strong prompt can contain both.

Context + Requirements

CHEESE · RECOMMENDED

3. Context and Input

Theory: The background and material the AI needs to know about.

Ask: What evidence or situation changes the answer?

Use only the approved policy text below.

TIP

Include what is relevant. More context is not always better context.

TOPPINGS · RECOMMENDED

4. Requirements and Details

Theory: The facts, limits and audience needs the answer must include.

Ask: What must be present for the result to work?

Explain acronyms. Keep it under 400 words. Include owners.

TIP

Write requirements so a person can check whether each one is present.

PRO TIP · KEEP SOURCE AND INSTRUCTIONS APART

Label source material clearly. This improves readability, but it does not prevent prompt injection.

Style + Format + Rules

SAUCE · OPTIONAL

5. Style and Quality

How it should sound or look, and how polished it needs to be.

Plain English, calm and direct.

BOTTOM BUN · RECOMMENDED

6. Output Format

The shape of the answer: list, table, code, image size and so on.

One title, five numbered steps, then a glossary.

WRAPPER · RECOMMENDED

7. Rules and Boundaries

What the AI must not do, change or include.

Do not invent policy details. Mark missing information.

COUNTER TIP

If Style genuinely does not matter, mark it “Not needed” and record why.

CHECK BEFORE USING

A stated boundary can still fail. Use workflow controls and human review for consequential work.

From rough request to strong brief

ROUGH REQUEST

Summarise our returns policy.

WHAT IS MISSING?

Purpose, audience, source, must-have details, output shape and a rule for missing information.



SEVEN LAYERS

Goal: Help new shop staff answer routine returns questions consistently.

Task: Summarise the approved policy and extract the process.

Context and Input: Use only the policy text inside `<source>` tags.

Requirements and Details: Cover eligibility, time limits, proof of purchase, exceptions and escalation.

Style and Quality: Plain English for first-week staff.

Output Format: Five numbered steps, then a short glossary.

Rules and Boundaries: Do not invent details. Write "Not stated" when the source is silent.

PRO TIP

The best assembled order can differ from the burger's learning order. With long source material, test whether repeating the Task after the source improves results.

Pick the smallest useful technique

Does a clear direct prompt work?

YES → Keep it simple.

NO ↓

Is the pattern hard to explain?

YES → Show examples.

NO ↓

Is the job too large to inspect?

YES → Decompose or chain.

NO ↓

Does the answer need outside facts or calculation?

YES → Ground it or use a tool.

NO ↓

Will software consume the result?

YES → Use structured output.

NO →
iterate
and
evaluate.

GOLDEN RULE

A technique is useful only when it improves the real result enough to justify its added effort, cost and review burden.

Direct prompting

Professional term: **Zero-shot prompting**

STARTER

Theory

Ask for a familiar task with clear instructions and no worked example.

Modern models often recognise common jobs such as summarising, rewriting and brainstorming. Starting simple gives you a useful baseline.

Best-fit use cases

- ✓ A short summary
- ✓ A first draft
- ✓ Common classifications
- ✓ Simple idea generation

Example

Sort each comment into Praise, Question or Complaint.
Return only the label and one quoted reason.

COUNTER TIP

Start here. Add a technique only when a real test shows the direct prompt is not enough.

PRO TIP

Keep the direct version as your baseline.
Compare every more complex version against the same test cases.

WATCH OUT

“Zero-shot” describes the prompt, not the quality. Unusual categories or ambiguous tasks may need examples.

Evidence [\[3\]](#) [\[4\]](#) · Test this technique on the model, version and task you actually use.

Role and audience

Professional term: **Role prompting**

STARTER

Theory

Give the AI a relevant point of view and name the people who will use the answer.

A role can focus vocabulary, priorities and tone. The audience tells the model what knowledge and language level to expect.

Best-fit use cases

- ✓ Teaching
- ✓ Accessibility review
- ✓ Audience-specific writing
- ✓ Simulated critique

Example

Act as a patient science tutor. Explain photosynthesis to a curious 10-year-old using one everyday comparison.

COUNTER TIP

Choose a role that changes the work. “Helpful expert” is usually too vague to add value.

PRO TIP

Specify responsibilities and review criteria, not invented credentials: “Review for plain language and explain every technical term.”

WATCH OUT

A role does not create expertise, evidence or accountability. “Act as a doctor” does not make medical advice reliable.

Evidence [2] · Test this technique on the model, version and task you actually use.

Show the pattern

Professional term: **One-shot and few-shot prompting**

SKILLED

Theory

Provide one or several examples that demonstrate the pattern you want the AI to follow.

Examples can make an unusual format, tone or decision rule easier to recognise than a long explanation.

Best-fit use cases

- ✓ Classification
- ✓ Brand voice
- ✓ Edge cases
- ✓ Specialised formats

Example

Comment: "Great mug, but checkout failed."

Labels: Product praise; Website problem

Now label: "Fast delivery, wrong colour."

COUNTER TIP

Include at least one awkward example. Easy examples rarely teach the difficult boundary.

PRO TIP

Keep examples relevant, varied and consistently formatted. Treat their labels like maintained test data.

WATCH OUT

Poor, biased or repetitive examples teach the wrong pattern. More examples also use more context.

Evidence [\[1\]](#) [\[2\]](#) [\[3\]](#) · Test this technique on the model, version and task you actually use.

Label the parts

Professional term: **Structured prompting and delimiters**

STARTER

Theory

Use headings, tags or fences to separate instructions, source material, examples and output rules.

Clear sections reduce accidental mixing and make a long prompt easier for people to read and maintain.

Best-fit use cases

- ✓ Long source documents
- ✓ Reusable templates
- ✓ Several inputs
- ✓ Document analysis

Example

TASK: Summarise the policy.
SOURCE: <policy>...</policy>
RULES: Use only the source.
FORMAT: Five bullets.

COUNTER TIP

Use names that describe the job: TASK, SOURCE, RULES and OUTPUT. Keep the labels consistent.

PRO TIP

Choose Markdown, XML or another stable convention. Structure for maintainability, not decoration.

WATCH OUT

A label is not a security wall. Malicious instructions inside a document can still influence a model.

Evidence [\[2\]](#) [\[3\]](#) [\[12\]](#) · Test this technique on the model, version and task you actually use.

Set a checkable finish line

Professional term: **Constraint prompting**

STARTER

Theory

State the limits and qualities that make an answer acceptable.

“Make it good” is hard to inspect. Concrete requirements make the target visible to both the AI and the reviewer.

Best-fit use cases

- ✓ Publishing
- ✓ Interface copy
- ✓ Regulated wording
- ✓ Consistent deliverables

Example

Give three options. Each must be under 40 words, use plain English and include one drawback.

COUNTER TIP

If rules compete, rank them. Say which requirement matters most.

PRO TIP

Turn important requirements into acceptance criteria that can be checked by code or a human reviewer.

WATCH OUT

Too many or conflicting constraints can reduce quality or cause silent trade-offs.

Evidence [2] [3] · Test this technique on the model, version and task you actually use.

Divide the work

Professional term: **Task decomposition and prompt chaining**

SKILLED

Theory

Split a complex job into smaller prompts whose results can be checked before the next stage.

Smaller stages are easier to inspect, repair and approve than one large request that hides several decisions.

Best-fit use cases

- ✓ Research
- ✓ Long documents
- ✓ Data cleaning
- ✓ Approval workflows

Example

Situation: A policy summary mixes extraction, verification and writing.

Basic prompt:

Summarise this policy accurately.

Technique-enabled prompt:

Run 1: Extract claims.

Run 2: Attach source evidence.

Run 3: Flag unsupported claims.

Run 4: Write only from verified claims.

What to inspect: Check each intermediate result before it becomes the next input.

Does not prove: Smaller stages do not guarantee that the evidence is complete.

COUNTER TIP

Give every stage a visible input and output.
Stop when a required check fails.

PRO TIP

Store intermediate results and evaluation notes. This creates a traceable workflow and makes regressions easier to find.

WATCH OUT

Errors can travel from one stage to the next. Chaining adds time, cost and places that can fail.

Evidence [2] · Test this technique on the model, version and task you actually use.

Support multi-step reasoning

Professional term: **Worked-reasoning examples**

PRO

Theory

For some difficult tasks, a worked example can demonstrate a useful method before a new problem.

Original chain-of-thought studies found gains on some arithmetic, commonsense and symbolic reasoning benchmarks.

Best-fit use cases

- ✓ Multi-step maths
- ✓ Logic exercises
- ✓ Planning with dependencies
- ✓ Rule-based decisions

Example

Situation: A scheduling problem has several dependencies and one valid result.

Basic prompt:

Create the schedule.

Technique-enabled prompt:

Here is one solved example with its final constraint checks. Use the same method for the new case. Return the schedule and a concise constraint checklist.

What to inspect: Verify every listed constraint against the final schedule.

Does not prove: A fluent reasoning trace does not prove the schedule is valid.

COUNTER TIP

Ask for concise checks or calculations you can verify, not a long performance of certainty.

PRO TIP

Reasoning behaviour differs by model. Current reasoning models may work best with simple outcome-level instructions.

WATCH OUT

A convincing explanation can support a wrong answer. Exposed reasoning is not proof, and some providers advise against asking for it.

Evidence [5] [13] · Test this technique on the model, version and task you actually use.

Compare independent attempts

Professional term: **Self-consistency**

PRO

Theory

Generate several independent attempts, compare their conclusions and verify the leading answer.

Multiple paths can reveal instability and, on some checkable reasoning tasks, improve benchmark results.

Best-fit use cases

- ✓ Arithmetic
- ✓ Multiple-choice reasoning
- ✓ Option comparison
- ✓ Uncertainty discovery

Example

Situation: A checkable reasoning problem gives unstable answers across runs.

Basic prompt:

Solve the problem once.

Technique-enabled prompt:

Run the same prompt separately several times with sampling enabled. Aggregate the final answers, then verify the leading answer with a calculator.

What to inspect: Record disagreement, the aggregation rule and the independent verification.

Does not prove: Agreement across sampled paths does not prove truth.

COUNTER TIP

Use disagreement as a signal to investigate.
Do not hide it.

PRO TIP

The original method samples diverse reasoning paths and aggregates answers. Implementation is model/API-dependent and may require separate calls.

WATCH OUT

Agreement is not truth. Several runs can repeat the same misconception or biased assumption.

Evidence [14] · Test this technique on the model, version and task you actually use.

Ground the answer

Professional term: **Retrieval-augmented generation (RAG)**

SKILLED

Theory

Supply relevant external material so the answer can use information beyond the model's built-in knowledge.

Retrieval can bring current or organisation-specific evidence into the prompt and make sources inspectable.

Best-fit use cases

- ✓ Policy questions
- ✓ Company handbooks
- ✓ Current information
- ✓ Large document collections

Example

Situation: A team needs answers from a current internal handbook.

Basic prompt:

What is our returns process?

Technique-enabled prompt:

Retrieve the three most relevant approved passages. Answer only from them. Cite the passage used for every policy claim and mark any gap "Not stated".

What to inspect: Open every citation and confirm it supports the nearby claim.

Does not prove: Good citations do not prove the retrieval system found every relevant passage.

COUNTER TIP

Open the cited passage yourself. A citation is useful only when it actually supports the claim.

PRO TIP

Evaluate retrieval and generation separately: first "Did we fetch the right evidence?", then "Did the answer use it correctly?"

WATCH OUT

Retrieval can return stale, irrelevant or malicious material. RAG does not automatically remove hallucination, injection or source-quality risks.

Evidence [\[10\]](#) [\[12\]](#) [\[15\]](#) · Test this technique on the model, version and task you actually use.

Use tools for evidence or action

Professional term: **Tool calling**; ReAct is one advanced pattern

PRO

Theory

Let the model call an approved tool for information, calculation or a permitted action, then use the result.

Tools can provide current data, exact calculations and controlled access to systems the model cannot know by itself.

Best-fit use cases

- ✓ Search
- ✓ Calculations
- ✓ Databases
- ✓ Approved business actions

Example

Situation: A price conversion needs a live rate and could trigger a payment.

Basic prompt:

Convert €500 to US dollars and pay the invoice.

Technique-enabled prompt:

Call the approved rate tool. Show the rate, source and timestamp. Calculate the total. Stop and request approval before any payment action.

What to inspect: Check the tool result, arguments, user permission and approval point.

Does not prove: A correct tool call does not authorise the next action.

COUNTER TIP

Use read-only tools first. Put approval immediately before any action that changes data or spends money.

PRO TIP

ReAct specifically interleaves reasoning, actions and observations. Ordinary function calling is broader. Validate names and arguments in code.

WATCH OUT

Prompt injection can redirect tool use. A prompt asking the model to behave safely is not an access-control system.

Evidence [\[12\]](#) [\[16\]](#) [\[19\]](#) · Test this technique on the model, version and task you actually use.

Give the answer a schema

Professional term: **Structured outputs**

PRO

Theory

Require named fields and data types so software can reliably locate each part of an answer.

On supported APIs, schema-constrained output is stronger than simply asking for “valid JSON”.

Best-fit use cases

- ✓ Data extraction
- ✓ Form filling
- ✓ UI content
- ✓ Application workflows

Example

Situation: Software must import order details without guessing where each value is.

Basic prompt:

Return the order as JSON.

Technique-enabled prompt:

Use the required schema: product (string), quantity (integer ≥ 1), dietary_warning (string or null). Reject extra fields.

What to inspect: Validate the structure, then separately check the values against the source.

Does not prove: Valid JSON does not prove that a product, quantity or warning is correct.

COUNTER TIP

Use a schema when software consumes the result. Use a table when a person consumes it.

PRO TIP

Validate meanings, ranges and permissions after parsing. Keep the schema small and version it like code.

WATCH OUT

Correct shape does not mean correct facts. Features and supported schema keywords differ by provider.

Evidence [\[6\]](#) [\[7\]](#) · Test this technique on the model, version and task you actually use.

Point to the evidence

Professional term: **Multimodal prompting**

SKILLED

Theory

Combine text with images, audio or files, and tell the model exactly which evidence matters.

“Look at this” is vague. Named inputs, regions and questions help the model focus on the intended evidence.

Best-fit use cases

- ✓ Charts
- ✓ Scanned forms
- ✓ Product photos
- ✓ Mixed document packs

Example

In image 2, read only the departure-time column. Return the city and time for rows marked Delayed.

COUNTER TIP

Number the files or images. Refer to them by number and name the detail to inspect.

PRO TIP

Test at realistic image sizes and qualities. Keep OCR or computer-vision checks for critical values.

WATCH OUT

Models can misread small text, counts and spatial relationships. Media can also contain hidden instructions.

Evidence **[9]** **[12]** · Test this technique on the model, version and task you actually use.

Brief a visual

Professional term: **Image prompting**

SKILLED

Theory

Describe the intended use, subject, scene, composition, style, lighting and constraints.

A visual brief gives the model relationships to aim for, while leaving room for variation.

Best-fit use cases

- ✓ Illustrations
- ✓ Campaign concepts
- ✓ Product scenes
- ✓ Image edits

Example

A4 cover illustration; seven-layer burger, front view; warm cream; tomato and mustard accents; blank space above; no text or logos.

COUNTER TIP

For an edit, say what changes and what must remain unchanged.

PRO TIP

Treat camera words as visual cues, not literal camera controls. Generate, inspect and revise one issue at a time.

WATCH OUT

Check text, anatomy, counts, likeness, rights and representation. Repeated edits can drift from the original.

Evidence **[8]** · Test this technique on the model, version and task you actually use.

Brief code with proof

Professional term: **Code prompting**

SKILLED

Theory

State the goal, environment, relevant files, constraints, expected behaviour and verification method.

Code depends on versions and surrounding systems. Tests turn “looks right” into something inspectable.

Best-fit use cases

- ✓ Bug fixes
- ✓ Refactoring
- ✓ Small features
- ✓ Test creation

Example

In this React/TypeScript app, add keyboard navigation without visual changes. Preserve the public API and run existing accessibility tests.

COUNTER TIP

Ask for the smallest safe change. Name what must not change.

PRO TIP

Provide acceptance criteria, sample inputs and expected outputs. Review the diff and run tests in the real environment.

WATCH OUT

Generated code may be insecure, incomplete or use invented packages and APIs. Verify dependencies before installing.

Evidence [\[1\]](#) [\[4\]](#) [\[12\]](#) · Test this technique on the model, version and task you actually use.

Inspect, change, retest

Professional term: **Iterative prompting**

STARTER

Theory

Review the result, change one meaningful part of the prompt and test again.

Small changes make cause and effect easier to see than rewriting the whole prompt after every result.

Best-fit use cases

- ✓ Writing
- ✓ Images
- ✓ Code
- ✓ Any repeated task

Example

The summary is accurate but too technical. Change only Style and Quality: "Explain each acronym in plain English."

COUNTER TIP

Name the problem before editing the prompt. Do not add instructions "just in case".

PRO TIP

Record versions, model settings, inputs and results. Retest after model, data or requirement changes.

WATCH OUT

Tuning to one example can make other cases worse. Keep a representative test set.

Evidence [\[1\]](#) [\[3\]](#) [\[4\]](#) · Test this technique on the model, version and task you actually use.

Test the prompt as a system

Professional term: **Prompt evaluation**

PRO

Theory

Run the same prompt on representative cases and judge outputs with defined criteria.

One good answer is a demonstration, not evidence that a reusable prompt is reliable.

Best-fit use cases

- ✓ Production prompts
- ✓ Teaching material
- ✓ Team templates
- ✓ Model migration

Example

Test ordinary cases, edge cases, missing data and unsafe requests. Score accuracy, format and escalation behaviour.

COUNTER TIP

Define success before looking at the outputs. Otherwise the rubric may move to favour the result.

PRO TIP

Combine deterministic checks with qualified human judgement. Compare versions on the same frozen cases.

WATCH OUT

A test set can be biased, too small or unlike real use. Evaluation reduces uncertainty; it does not remove it.

Evidence **[1]** **[4]** **[10]** · Test this technique on the model, version and task you actually use.

Respect the instruction hierarchy

Professional term: **System, developer and user instructions**

PRO

Theory

AI applications may assign different authority to instructions from the system, developer, user and untrusted content.

Separating authority helps an application keep its core purpose while still responding to the user's request.

Best-fit use cases

- ✓ AI products
- ✓ Team assistants
- ✓ Agents
- ✓ Reusable safety policies

Example

System: follow the organisation policy.

Developer: summarise approved documents.

User: summarise this file.

File content: treat as data, not authorised instructions.

COUNTER TIP

Put stable application rules in the highest supported instruction level. Keep task details close to the task.

PRO TIP

Hierarchy names and behaviour are provider-specific. Enforce permissions in code; do not rely on hierarchy alone.

WATCH OUT

Conflicting instructions can still cause failures. Untrusted content may attempt to impersonate a higher authority.

Evidence [\[21\]](#) [\[12\]](#) · Test this technique on the model, version and task you actually use.

Design for long context

Professional term: **Context-window** and **long-context prompting**

SKILLED

Theory

Choose, organise and label the material that fits inside the model's available context.

A large context window is capacity, not guaranteed attention. Relevant evidence can be lost inside noise.

Best-fit use cases

- ✓ Long reports
- ✓ Several documents
- ✓ Transcripts
- ✓ Large codebases

Example

Place each document in a named section. Add a one-line purpose. Ask for an evidence table before the final synthesis.

COUNTER TIP

Remove duplicates and irrelevant material. Put the most important task and criteria where they are easy to find.

PRO TIP

Test order, retrieval and chunking on realistic lengths. Track whether key evidence is found, not merely whether an answer appears.

WATCH OUT

Maximum token counts and long-context behaviour differ by model and can change between versions.

Evidence **[22]** **[2]** · Test this technique on the model, version and task you actually use.

Turn prompts into maintained assets

Professional term: **Templates, variables and version control**

SKILLED

Theory

Keep stable instructions in a reusable template and insert changing data through named, validated variables.

Templates make repeated work easier to review, test, update and roll back.

Best-fit use cases

- ✓ Team workflows
- ✓ Products
- ✓ Repeated reports
- ✓ Prompt libraries

Example

Template: "Summarise <source> for <audience> using <format>." Validate every variable before building the final prompt.

COUNTER TIP

Name variables by meaning, not position:
audience is clearer than input_2.

PRO TIP

Store prompts with code, review changes, keep test fixtures and tag releases. Record the model and settings used in evaluation.

WATCH OUT

A template can spread a mistake at scale. Variables may carry secrets or injection attempts.

Evidence **[1]** **[4]** · Test this technique on the model, version and task you actually use.

Explore alternative paths

Professional term: **Tree-of-Thought-style workflows**

PRO

Theory

Generate several candidate approaches, evaluate them and continue only with the most promising options.

Some planning and search problems benefit from considering alternatives before committing to one route.

Best-fit use cases

- ✓ Planning
- ✓ Puzzles
- ✓ Design options
- ✓ Decision preparation

Example

Propose three launch plans. Score each against budget, time and risk. Expand the top two, then compare their failure modes.

COUNTER TIP

Make the criteria visible. “Best” is meaningless until you say what matters.

PRO TIP

Tree-of-Thought is a research family, not a magic phrase. Implement branching and selection explicitly when the task justifies it.

WATCH OUT

More branches cost more and can create false confidence. The evaluator may share the generator’s blind spots.

Evidence **[20]** · Test this technique on the model, version and task you actually use.

Tune model settings carefully

Professional term: **Sampling and provider settings**

PRO

Theory

Some tools expose settings that influence variation, length, reasoning effort or reproducibility.

Settings can change the balance between consistency, exploration, cost and latency.

Best-fit use cases

- ✓ Creative variation
- ✓ Repeatable testing
- ✓ Reasoning tasks
- ✓ Production optimisation

Example

Keep the prompt and test set fixed. Compare the supported default with one lower-variation setting; record quality, time and cost.

COUNTER TIP

Change one setting at a time. Use the provider's current documentation for the exact model.

PRO TIP

A seed may improve repeatability without guaranteeing identical output. Some current reasoning models ignore or restrict familiar controls.

WATCH OUT

Parameter names and effects differ by provider and model. Old advice can become wrong after a model update.

Evidence [\[1\]](#) [\[4\]](#) [\[5\]](#) · Test this technique on the model, version and task you actually use.

Prompt across languages

Professional term: **Multilingual and translation prompting**

SKILLED

Theory

Name the source language, target language, audience, locale and terms that must remain consistent.

Translation quality depends on context, register and regional use—not only word substitution.

Best-fit use cases

- ✓ Translation
- ✓ Localisation
- ✓ Multilingual support
- ✓ Plain-language adaptation

Example

Translate German onboarding text into UK English for new retail staff. Preserve product names. Flag idioms with no direct equivalent.

COUNTER TIP

Ask a fluent speaker to review high-impact or public text. Include an approved glossary for repeated terms.

PRO TIP

Evaluate meaning, tone and locale separately. Back-translation can reveal differences but is not proof of correctness.

WATCH OUT

Low-resource languages, dialects and cultural references may be handled unevenly. Personal data rules still apply.

Evidence **[3]** **[10]** · Test this technique on the model, version and task you actually use.

Budget the workflow

Professional term: **Cost and latency trade-offs**

PRO

Theory

Every added example, retrieved document, model call and review step uses time and computing resources.

A small quality gain may not justify a slower, more expensive workflow for every request.

Best-fit use cases

- ✓ Production systems
- ✓ Batch work
- ✓ Repeated evaluation
- ✓ Agent workflows

Example

Compare direct prompting with a five-run method on the same cases. Record accuracy, median response time and cost per accepted result.

COUNTER TIP

Measure cost per useful, reviewed result—not cost per model call.

PRO TIP

Route simple tasks to the simplest proven workflow. Reserve retrieval, multiple samples and high reasoning effort for cases that need them.

WATCH OUT

Optimising only for speed or cost can hide quality and safety failures. Optimising only for quality can make a workflow impractical.

Evidence [\[1\]](#) [\[4\]](#) [\[14\]](#) · Test this technique on the model, version and task you actually use.

Check rights and provenance

Professional term: Copyright, licensing, likeness and content provenance

SKILLED

Theory

Record where source material and generated assets came from, what permissions apply and what claims remain uncertain.

A technically good output may still be unsuitable to publish, reuse or attribute.

Best-fit use cases

- ✓ Publications
- ✓ Marketing
- ✓ Image generation
- ✓ Training and reference material

Example

Use only organisation-owned or licensed source images. Record the licence, creator, generation tool and review decision with the final asset.

COUNTER TIP

Do not paste private reference material into a tool until its use is authorised.

PRO TIP

Keep a provenance record and review the current law and platform terms for the relevant jurisdiction and use.

WATCH OUT

Legal status varies by country and facts. A prompt cannot grant rights or prove consent.

Evidence [23] [10] · Test this technique on the model, version and task you actually use.

BITE the prompt before use

BITE checks whether the prompt says what it needs to say. It is a reflection tool —not a score and not a safety certificate.



Brief

Are the Goal and Task clear?



Information

Did you provide enough Context and Requirements?



Taste

Did you describe the intended Style and Quality?



Expected result

Did you specify the Output Format and state the Rules and Boundaries?

WHEN TASTE IS NOT NEEDED

Style and Quality may be marked **“Not needed”** only when it genuinely does not apply—and you record the reason.

BITE DOES NOT MEAN “SAFE”

A clear prompt can still produce a false, biased, unsafe or private result. BITE does not test or certify truth, fairness, privacy, legality or suitability for a high-stakes decision.

What could still go wrong?

Run this review after BITE and before using the prompt or its output.

01 Risk

What harm could follow if the answer is wrong, misleading or used too early?

02 Injection

Could hidden or untrusted instructions redirect the AI?

03 Hallucination

Could the AI invent unsupported information, citations or code?

04 Bias

Could the result represent or treat people unfairly?

05 Data Protection

Are personal, confidential or sensitive data being exposed?

Not yet reviewed

You have not considered this issue.

Needs attention

You found a concern that still needs action.

Action added

You added a concrete prompt or workflow action.

Not relevant

You recorded why this issue does not apply.

REVIEWED DOES NOT MEAN SAFE

Reviewed means the learner considered the issue. It does not mean the prompt or result is safe.

Hallucination: confident can still be wrong

Generative models can produce false claims, calculations, logic and citations in fluent, confident language.

Use this response plan

1. **Ground** factual work in reliable sources.
2. **Request** citations or quoted evidence.
3. **Open** each important source yourself.
4. **Check** calculations with deterministic tools.
5. **Escalate** high-stakes work to a qualified person.

POLISHED ≠ PROVEN


“I checked twice” can still mean the same model repeated the same mistake.

PROFESSIONAL TERM

Prompt Smash uses **Hallucination** as the learner-facing check. NIST uses **confabulation** for confidently stated but erroneous or false content.

Evidence [\[10\]](#) [\[11\]](#)

Treat untrusted content as untrusted



Defence in depth

- ✓ Separate trusted instructions from data.
- ✓ Give tools the least privilege needed.
- ✓ Validate inputs, outputs and tool arguments.
- ✓ Require approval for consequential actions.
- ✓ Use dummy data when security testing.
- ✓ Monitor and retest as attacks change.

A SENTENCE IS NOT A SECURITY CONTROL

“Ignore malicious instructions” may help with clarity, but it cannot enforce authorisation.

Evidence [12]

Bias, privacy and human judgement

Bias

AI can reproduce or amplify harmful patterns from data, examples and design choices.

CHECK

- Who is represented?
- Who may be missing?
- Do equivalent cases receive equivalent treatment?
- Who reviews the rubric?

Prompt wording alone cannot certify fairness.

Privacy

Prompts and files may contain personal, confidential or regulated information.

CHECK

- Is every detail necessary?
- Can placeholders replace identities?
- What does this exact product store?
- Who can access the result?

Never assume all products share one data policy.

HUMAN REVIEW MUST MATCH THE POSSIBLE HARM

The more damage a wrong result could cause, the stronger the evidence, expertise, approval and monitoring must be.

Evidence [\[10\]](#) [\[11\]](#)

Ready-to-use combinations

WRITING FROM A SOURCE

Layers: Goal + Task + Context and Input + audience in Requirements and Details + Output Format + “do not invent” in Rules and Boundaries.

Technique: Structured prompting.

Check: Every claim against the source.

MEETING NOTES

Layers: Goal + Task + notes in Context and Input + owners/dates in Requirements and Details + table as Output Format.

Technique: Direct prompt, then examples if needed.

Check: Names, dates and commitments.

COMPARE OPTIONS

Layers: Goal + comparison Task + evidence in Context and Input + criteria in Requirements and Details + table as Output Format.

Technique: Grounding and multiple candidates.

Check: Source quality and assumptions.

CLASSIFY FEEDBACK

Layers: Goal + classification Task + comments in Context and Input + multi-label rule + JSON or table as Output Format.

Technique: Few-shot and evaluation.

Check: Edge cases and group differences.

WEBSITE IMAGE

Layers: use as Goal + generate/edit Task + references in Context and Input + composition + Style and Quality + size as Output Format + preservation Rules and Boundaries.

Technique: Image briefing and iteration.

Check: text, rights and crop.


SMALL CODE FEATURE

Layers: outcome as Goal + implementation Task + repository Context and Input + acceptance criteria + diff as Output Format + no-secrets Rules and Boundaries.

Technique: Decomposition and tests.

Check: tests, diff and security.

The complete workflow



For one-off work

Use a small sample, inspect the answer and keep a person responsible for the final use.

For repeated work

Version the prompt, freeze representative test cases, define metrics and monitor failures over time.

For high-stakes work

Use qualified review, controlled data access, documented approval and a way to stop or correct the system.

ACCESSIBLE OUTPUTS

When the result will become digital content, include accessibility requirements and verify the finished work against the relevant standard. [\[18\]](#)

REMEMBER

A better prompt can improve an AI interaction. It cannot turn a probabilistic model into a guaranteed source of truth.

Pocket glossary

AI model

A system that finds patterns in data and generates a response.

Prompt

The instructions and information given to a model.

Prompt design

Writing a clear prompt for a task.

Prompt engineering

Systematically building, testing and maintaining prompts and workflows.

Zero-shot

Asking without a worked example.

One-shot / few-shot

Showing one or several examples.

Grounding

Connecting an answer to supplied evidence.

RAG

Retrieving relevant material and adding it to the model's context.

Schema

A formal definition of fields and data types.

Multimodal

Using more than one kind of input, such as text, images, audio or files.

Hallucination

A false or unsupported response that may sound confident. NIST uses the related term "confabulation".

Prompt injection

Untrusted instructions that try to redirect an AI system.

Evaluation

Testing outputs against defined criteria and representative cases.

Deterministic check

A check that gives the same result for the same input, such as a schema validator.

Human in the loop

A person reviews or approves a decision before it is used.

Least privilege

Giving a system only the access it needs.

Research notes and sources

This guide combines Prompt Smash's approved method with current provider guidance, original research and risk-management sources.

Provider advice can change. Links below open the original source and were reviewed on 8 October 2026.

1. **OpenAI, "Prompting" and "Prompt engineering"**
[Open the original source ↗](#) · reviewed 8 October 2026
2. **Anthropic, "Prompting best practices"**
[Open the original source ↗](#) · reviewed 8 October 2026
3. **Google, "Prompt design strategies"**
[Open the original source ↗](#) · reviewed 8 October 2026
4. **OpenAI, "Model optimization"**
[Open the original source ↗](#) · reviewed 8 October 2026
5. **OpenAI, "Reasoning best practices"**
[Open the original source ↗](#) · reviewed 8 October 2026
6. **OpenAI, "Structured model outputs"**
[Open the original source ↗](#) · reviewed 8 October 2026
7. **Google, "Structured outputs"**
[Open the original source ↗](#) · reviewed 8 October 2026
8. **OpenAI, "Image prompting"**
[Open the original source ↗](#) · reviewed 8 October 2026
9. **Google, "Files API"**
[Open the original source ↗](#) · reviewed 8 October 2026
10. **NIST, Artificial Intelligence Risk Management Framework: Generative AI Profile**
[Open the original source ↗](#) · reviewed 8 October 2026
11. **NIST, AI Risk Management Framework**
[Open the original source ↗](#) · reviewed 8 October 2026
12. **OWASP, LLM Prompt Injection Prevention Cheat Sheet**
[Open the original source ↗](#) · reviewed 8 October 2026

Sources continued

Original research and advanced workflow sources.

13. **Wei et al., “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models”**
[Open the original source ↗](#) · reviewed 8 October 2026
14. **Wang et al., “Self-Consistency Improves Chain of Thought Reasoning in Language Models”**
[Open the original source ↗](#) · reviewed 8 October 2026
15. **Lewis et al., “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks”**
[Open the original source ↗](#) · reviewed 8 October 2026
16. **Yao et al., “ReAct: Synergizing Reasoning and Acting in Language Models”**
[Open the original source ↗](#) · reviewed 8 October 2026
17. **Schulhoff et al., “The Prompt Report: A Systematic Survey of Prompting Techniques”**
[Open the original source ↗](#) · reviewed 8 October 2026
18. **W3C, Web Content Accessibility Guidelines (WCAG) 2.2**
[Open the original source ↗](#) · reviewed 8 October 2026
19. **OpenAI, “Function calling”**
[Open the original source ↗](#) · reviewed 8 October 2026
20. **Yao et al., “Tree of Thoughts: Deliberate Problem Solving with Large Language Models”**
[Open the original source ↗](#) · reviewed 8 October 2026
21. **OpenAI, “The Instruction Hierarchy”**
[Open the original source ↗](#) · reviewed 8 October 2026
22. **Google, “Long context”**
[Open the original source ↗](#) · reviewed 8 October 2026
23. **U.S. Copyright Office, Copyright and Artificial Intelligence**
[Open the original source ↗](#) · reviewed 8 October 2026

HOW TO USE THESE SOURCES

Provider documentation describes current product behaviour. Research papers describe tested methods under specific conditions. Neither guarantees the same result on every model or task.

Your final tray check



Seven layers: Goal → Task → Context and Input → Requirements and Details → Style and Quality → Output Format → Rules and Boundaries.

Before you prompt

- ✓ Do I know the Goal and Task?
- ✓ Is the input relevant and safe to share?
- ✓ Can I describe a checkable result?

Before you use the output

- ✓ Did I run BITE?
- ✓ Did I review Risk, Injection, Hallucination, Bias and Data Protection?
- ✓ Did I verify the evidence?
- ✓ Does a qualified person need to approve it?

ONE USEFUL HABIT

Save the prompt, the input, the output and the review decision together when the work matters.

Build clearly. Test honestly. Review responsibly.